

Darts

Max Caragozian

April 21, 2025

The tap-room fire is alight again
and a new stock of darts laid in;
serviceable and well-feathered,
that fly true and will see us into
another spring.

“A Chiltern Autumn” [4]

1 Introductction

When I moved to Rochester this past January, I bought a dartboard and hung it in the basement of my college house. It ended up being a worthwhile purchase, both as a way to get to know my new roommates and as time killer for what ended up being a long, snowy winter. I started reading into the mathematics of darts and came across the paper “A Statistician Plays Darts” [2], which quantifies how a dart thrower’s skill affects where he/she should aim.

In the simplest model, the authors assume that a person throws with a symmetrical bivariate normal distribution. They estimate the function $\mu^*(\sigma)$ that takes in a standard deviation σ and outputs a point μ that maximizes the expected score of a dart thrown with distribution $\mathcal{N}(\mu, \sigma^2)$.

Bad sentence -;. Unsurprisingly, the authors found that a person with a very small standard deviation should aim for the triple 20 (T20), as it has the highest point value on the board. The optimal aiming location stays in the T20 until approximately $\sigma = 16.9\text{mm}$, where the optimal location jumps to the T19, before eventually moving towards the center of the board.

At the outset of this project, my aim was to replicate some of the results of “A Statistician Plays Darts.” [2] Along the way, I found that there is more than one way to get there and learned much about global optimization algorithms. This paper documents the several methods I tried, how I implemented them, and how each method 111




Figure 1: Standard dartboard with point values [3]

2 Integration Brute Force

Before trying to find $\mu^*(\sigma)$, I wrote code that would find the expected score for someone throwing a dart with a given random distribution. We define a *sector* as a contiguous area of the dartboard with a single point value. Examples are the triple 20, the double bullseye, or the portion of single 19 that lies between the outer bullseye ring and the inner triple ring. Let S be the set of all sectors and X be a distribution with PDF f_x . To find the expected score of a distribution we compute:

$$E[\text{score}(X)] = \sum_{s \in S} \iint_s \text{score}_s \cdot f_x(x, y) dx dy. \quad (1)$$

I used SciPy's *nquad* function to perform the integration. Originally, I used the *dblquad* method, which is specific to double integral, but I ran into a rounding issues when f_x was very small. The *nquad* function allows finer control and fixed the issue.

With the code implementation of Equation 1, I could fix σ and calculate the expected score for normal distributions centered at many different points μ across the dartboard, returning both the approximate global maximum (μ^*) and data that could be plotted as a heatmap.

For each pixel in the heat map I computed 82 double integrals: one for each sector of the dartboard. This was painfully slow, and I spent a lot of time trying to optimize the code by testing the value of the PDF at certain values across the sector so I could skip integrating over sectors that would contribute very little to the final expected score. After a lot of time tinkering with the




Figure 2: Heatmap for $\sigma = 26.9$. The green $+$ is the location of μ^*

threshold for rejecting a sector and how to distribute test points, etc., I came to the conclusion it was easier just to compute every integral. The speedup from my optimization was minimal, and I couldn't escape the fear that by skipping certain integrals I might materially affect the final calculation.

STUFF ABOUT TIMKNG

3 Global Optimization Algorithms

I looked for faster ways to compute μ^* for a fixed σ . We can think about the problem as maximizing the function

$$F(\mu|\sigma) = E[\text{score}(\mathcal{N}(\mu, \sigma^2))] \quad (2)$$

over all points on the dartboard μ . Because optimization algorithms by convention minimize functions rather than maximizing them, we end up computing

$$\min(-F(\mu|\sigma)). \quad (3)$$

SciPy has bindings for several different global minimization algorithms. I ended up settling on a stochastic algorithm called *basin-hopping*. The basin-hopping algorithm takes in an initial guess, and then modifies it by a step of a

random size in a random direction. After landing at new point, the algorithm employs a local minimization technique to move “downhill” to a local minimum. It then compares the new local minimum with the most recent point. If the new minimum is lower than the previous one, it is always accepted. If it is not lower, there is a random process to decide whether or not to accept it, with lower values being more likely to be accepted. The process then repeats, with another random perturbation, a local minimization, and a comparison with the most recently accepted point. Over time, random jumps get smaller, and the algorithm ends when a point has remained unchallenged for some number of cycles. The distribution of jump sizes, the local minimization algorithm, and the acceptance criteria for new points can all be controlled by parameters [1].

We should note that basin-hopping is not guaranteed to find the global minimum. If the random jumps never find the right “basin”, the algorithm will return the wrong answer.

SciPy allows the programmer to pass callbacks to both the global and local minimizer, meaning one can record the progress of the algorithm. I wrote code to write the progress to a file and then plot the data over a heatmap such as that in Figure 2.

While the dartboard problem proved a good way to illustrate basin-hopping, basin-hopping was not a good way to solve my dartboard problem. There were two main issues, the of which first I have already noted: basin hopping is not guaranteed to find the true global minimum. The run in Figure 3 is an example of a time when basin-hopping failed to find the true global minimum. The other problem is that basin-stopping still relies on the slow process of integration for each point that it tests. To be sure, it is still faster than brute force. The basin-hopping run in Figure 3 needed only 56 evaluations of $F(\mu|\sigma)$, compared to 9000 for a 300×300 heatmap. But it was still slow and it got the wrong answer, so I decided to go back to square one.

4 Don’t Let 10 Minutes of Reading Save You From 10 Hours of Coding

References

- [1] The Scipy Community. *scipy.optimize.basinhopping*. URL: <https://docs.scipy.org/doc/scipy/reference/generated/scipy.optimize.basinhopping.html#scipy.optimize.basinhopping>.
- [2] Andrew Price Ryan Tibshirani and Jonathan Taylor. “A Statistician Plays Darts”. In: *Journal of the Royal Statistical Society Series A* 174 (2011), pp. 213–226.
- [3] Tijmen Stam. *Dartboard diagram.svg*. 2014. URL: https://commons.wikimedia.org/wiki/File:Dartboard_diagram.svg.

